

DS4001 Databases (7.5 credits)

# Lecture II – Query Optimization

Yuantao Fan  
yuantao.fan@hh.se

Halmstad University

# Introduction

- SQL queries are declarative...
  - You wrote what you want
  - Not how to process data in tables to acquire the result
  - Difference in performance depending on the methods/plans
- First query optimizer – IBM system R
  - Human can write better queries compared to DBMS
- Query optimization
  - Heuristics - improve the query by rewriting inefficient operations, without evaluating the cost
  - Cost-based search
    - Generate a set of equivalent plans, compute their cost and select the most efficient one

# Example – Predicate Pushdown

```
SELECT t.full_name, tt.cid  
FROM Teacher AS t, Teaches AS tt  
WHERE t.tid = tt.tid  
      AND tt.hours > 50
```

$\pi_{\text{full\_name, cid}}(\sigma_{\text{hours} > 50}(\text{Teacher} \bowtie \text{Teaches}))$

Teacher(tid, full\_name, age, nationality)

| tid | full_name     | age | nationality |
|-----|---------------|-----|-------------|
| 11  | John Smith    | 42  | America     |
| 22  | Jens Jonathon | 31  | Sweden      |
| 33  | Stefan Miller | 39  | Sweden      |
| 44  | Kayle Persson | 33  | UK          |

Teaches(tid, cid, hours)


| <u>tid</u> | <u>cid</u> | hours |
|------------|------------|-------|
| 11         | 1          | 80    |
| 11         | 2          | 100   |
| 22         | 4          | 50    |
| 33         | 4          | 50    |
| 44         | 3          | 100   |

# Query Plan

```
SELECT t.full_name, tt.cid  
FROM Teacher AS t, Teaches AS tt  
WHERE t.tid = tt.tid  
      AND tt.hours > 50
```



- A list of instructions that the database needs to follow in order to execute a query on the data
- Query tree is a tree data structure that corresponds to an extended relational algebra expression
- The operations are arranged in a tree and the input data (i.e. tables/relations) flow from the bottom (leaves) to the top (root)

$\pi_{\text{full\_name, cid}}(\sigma_{\text{hours} > 50}(\text{Teacher} \bowtie \text{Teaches}))$



# Example – Predicate Pushdown

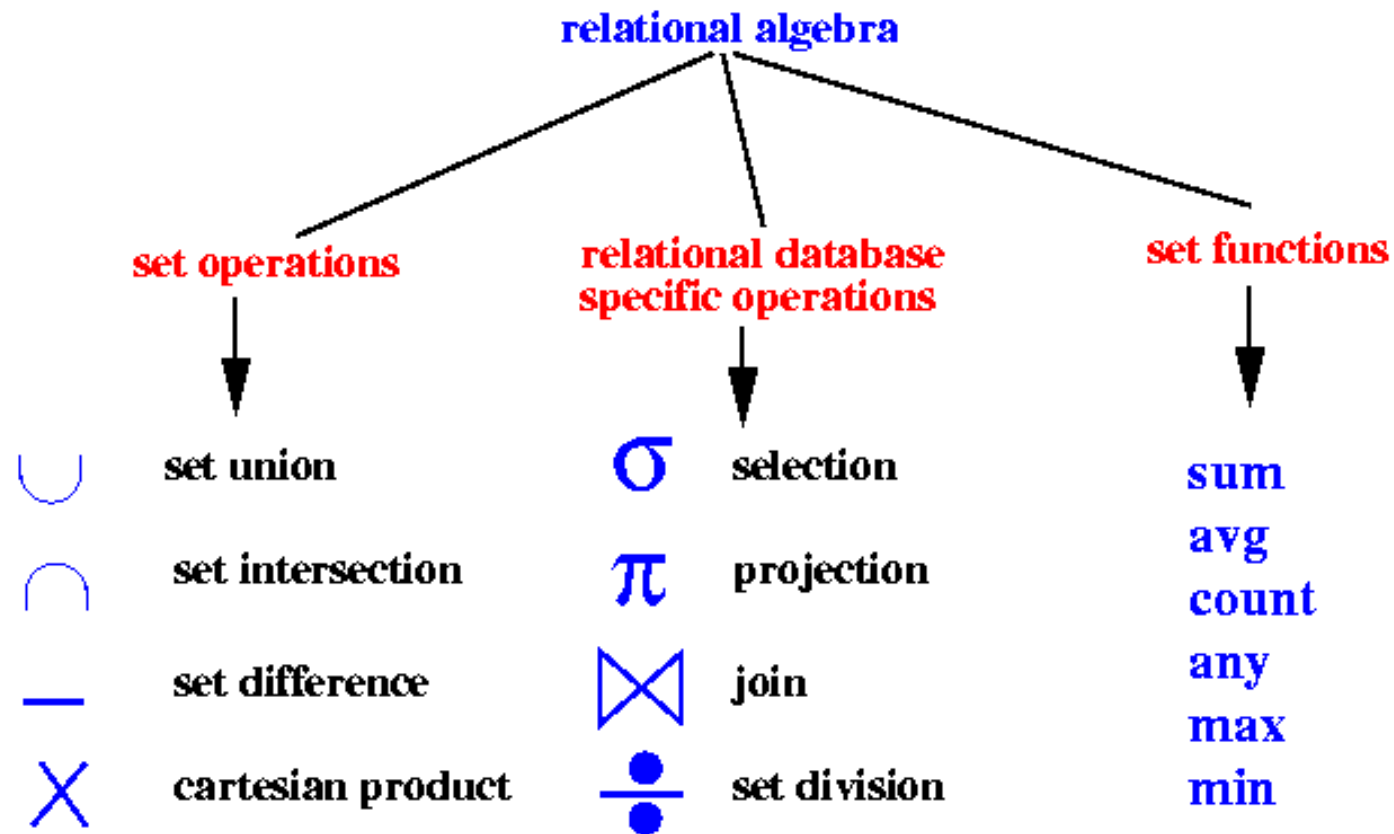
```
SELECT t.full_name, tt.cid
FROM Teacher AS t, Teaches AS tt
WHERE t.tid = tt.tid
      AND tt.hours > 50
```



$$\pi_{full\_name, cid}(\sigma_{hours>50}(Teacher \bowtie Teaches)) = \pi_{full\_name, cid}(Teacher \bowtie \sigma_{hours>50}(Teaches))$$

# Recall Relational Algebra

- Fundamental operations to retrieve and manipulate tuples in a relation



# Relational Algebra - Selection

- Syntax
  - $\sigma_{predicate}(R)$
- Choose a subset of the tuples from a relation that satisfies a selection predicate
  - Filtering using predicates
  - Combine multiple predicates

$\sigma_{Tid='a1'}(T)$

| Tid | Cid |
|-----|-----|
| a1  | 10  |

```
SELECT *  
FROM T  
WHERE Tid='a1';
```

T(Tid, Cid)

| Tid | Cid |
|-----|-----|
| a1  | 10  |
| a2  | 11  |
| a3  | 12  |
| a3  | 13  |

$\sigma_{Tid='a3' \wedge Cid='12'}(T)$

| Tid | Cid |
|-----|-----|
| a3  | 12  |

```
SELECT *  
FROM T  
WHERE Tid='a3' AND Cid='12';
```

# Relational Algebra - Projection

- Syntax
  - $\pi_{A1,A2,\dots,An}(R)$
- Generate a relation with tuples that contains attributes specified
  - Can manipulate the values
  - Can rearrange the order

T(Tid, Cid)

| Tid | Cid |
|-----|-----|
| a1  | 10  |
| a2  | 11  |
| a3  | 12  |
| a3  | 13  |

$\pi_{Tid,Cid*2}(\sigma_{Tid='a2'}(T))$

| Tid | Cid |
|-----|-----|
| a2  | 24  |

```
SELECT Tid, Cid*2
FROM T
WHERE Tid='a2';
```



# Relational Algebra - Union

- Syntax
  - $(A \cup B)$
- Collect all tuples in both relations

A(Tid, Cid)

| Tid | Cid |
|-----|-----|
| a1  | 10  |
| a2  | 11  |

B(Tid, Cid)

| Tid | Cid |
|-----|-----|
| a2  | 11  |
| a3  | 13  |

$(A \cup B)$

| Tid | Cid |
|-----|-----|
| a1  | 10  |
| a2  | 11  |
| a2  | 11  |
| a3  | 13  |

```
(SELECT * FROM A)
UNION ALL
(SELECT * FROM B);
```

# Relational Algebra - Intersection

- Syntax
  - $(A \cap B)$
- Collect tuples that appear in both relations

A(Tid, Cid)

| Tid | Cid |
|-----|-----|
| a1  | 10  |
| a2  | 11  |

B(Tid, Cid)

| Tid | Cid |
|-----|-----|
| a2  | 11  |
| a3  | 13  |

$(A \cap B)$

| Tid | Cid |
|-----|-----|
| a2  | 11  |

```
(SELECT * FROM A)
INTERSECT
(SELECT * FROM B);
```

# Relational Algebra - Product

- Syntax
  - $(A \times B)$
- Generate a relation that contains all possible combinations of the tuples from both relations

`SELECT * FROM A, B;`

`SELECT * FROM A CROSS JOIN B;`

A(Tid, Cid)

| Tid | Cid |
|-----|-----|
| a1  | 10  |
| a2  | 11  |

B(Tid, Cid)

| Tid | Cid |
|-----|-----|
| a2  | 11  |
| a3  | 13  |

$(A \times B)$

| A.Tid | A.Cid | B.Tid | B.Cid |
|-------|-------|-------|-------|
| a1    | 10    | a2    | 11    |
| a1    | 10    | a3    | 13    |
| a2    | 11    | a2    | 11    |
| a2    | 11    | a3    | 13    |

# Relational Algebra - Join

- Syntax
  - $(A \bowtie B)$
- $R \bowtie_{\langle \text{join condition} \rangle} S$
- Generate a relation that contains all tuples with a common value(s) of one (or more) attribute(s)

A(Tid, Cid)

| Tid | Cid |
|-----|-----|
| a1  | 10  |
| a2  | 11  |

B(Tid, Cid)

| Tid | Cid |
|-----|-----|
| a2  | 11  |
| a3  | 13  |

$(A \bowtie B)$

| Tid | Cid |
|-----|-----|
| a2  | 11  |

SELECT \* FROM A NATURAL JOIN B;

# The Division operation $\div$

- Syntax
  - $R \div S$
- Produces a relation  $R(X)$  that includes all tuples  $t[X]$  in  $S(Z)$  that appear in  $R$  in combination with every tuple from  $S(Y)$ , where  $Z = X \cup Y$
- “An example is Retrieve the names of employees who work on all the projects that ‘John Smith’ works on.”

**Figure 8.8**

The DIVISION operation. (a) Dividing SSN\_PNOS by SMITH\_PNOS. (b)  $T \leftarrow R \div S$ .

(a)

SSN\_PNOS

| Essn      | Pno |
|-----------|-----|
| 123456789 | 1   |
| 123456789 | 2   |
| 666884444 | 3   |
| 453453453 | 1   |
| 453453453 | 2   |
| 333445555 | 2   |
| 333445555 | 3   |
| 333445555 | 10  |
| 333445555 | 20  |
| 999887777 | 30  |
| 999887777 | 10  |
| 987987987 | 10  |
| 987987987 | 30  |
| 987654321 | 30  |
| 987654321 | 20  |
| 888665555 | 20  |

SMITH\_PNOS

| Pno |
|-----|
| 1   |
| 2   |

SSNS

| Ssn       |
|-----------|
| 123456789 |
| 453453453 |

(b)

R

| A  | B  |
|----|----|
| a1 | b1 |
| a2 | b1 |
| a3 | b1 |
| a4 | b1 |
| a1 | b2 |
| a3 | b2 |
| a2 | b3 |
| a3 | b3 |
| a4 | b3 |
| a1 | b4 |
| a2 | b4 |
| a3 | b4 |

S

| A  |
|----|
| a1 |
| a2 |
| a3 |

T

| B  |
|----|
| b1 |
| b4 |

Textbook p.256

# Aggregate Functions and Grouping

- Syntax
  - $\langle \text{grouping attribute} \rangle \mathfrak{F} \langle \text{function list} \rangle (R)$

Where

- $\langle \text{grouping attributes} \rangle$  is a list of attributes of the relation specified in  $R$ ; and
- $\langle \text{function list} \rangle$  is a list of ( $\langle \text{function} \rangle$   $\langle \text{attribute} \rangle$ ) pairs. In each such pair;
- $\langle \text{function} \rangle$  is one of the allowed functions—such as SUM, AVERAGE, MAXIMUM, MINIMUM, COUNT—and  $\langle \text{attribute} \rangle$  is an attribute of the relation specified by  $R$ .

$\rho_R(Dno, No\_of\_employees, Average\_sal) (Dno \mathfrak{F} COUNT Ssn, AVERAGE Salary (EMPLOYEE))$

(a)

| Dno | No_of_employees | Average_sal |
|-----|-----------------|-------------|
| 5   | 4               | 33250       |
| 4   | 3               | 31000       |
| 1   | 1               | 55000       |

(b)

| Dno | Count_ssn | Average_salary |
|-----|-----------|----------------|
| 5   | 4         | 33250          |
| 4   | 3         | 31000          |
| 1   | 1         | 55000          |

(c)

| Count_ssn | Average_salary |
|-----------|----------------|
| 8         | 35125          |

**Figure 8.10**

The aggregate function operation.

- $\rho_R(Dno, No\_of\_employees, Average\_sal) (Dno \mathfrak{F} COUNT Ssn, AVERAGE Salary (EMPLOYEE)).$
- $Dno \mathfrak{F} COUNT Ssn, AVERAGE Salary (EMPLOYEE).$
- $\mathfrak{F} COUNT Ssn, AVERAGE Salary (EMPLOYEE).$

# Relational Algebra and operations

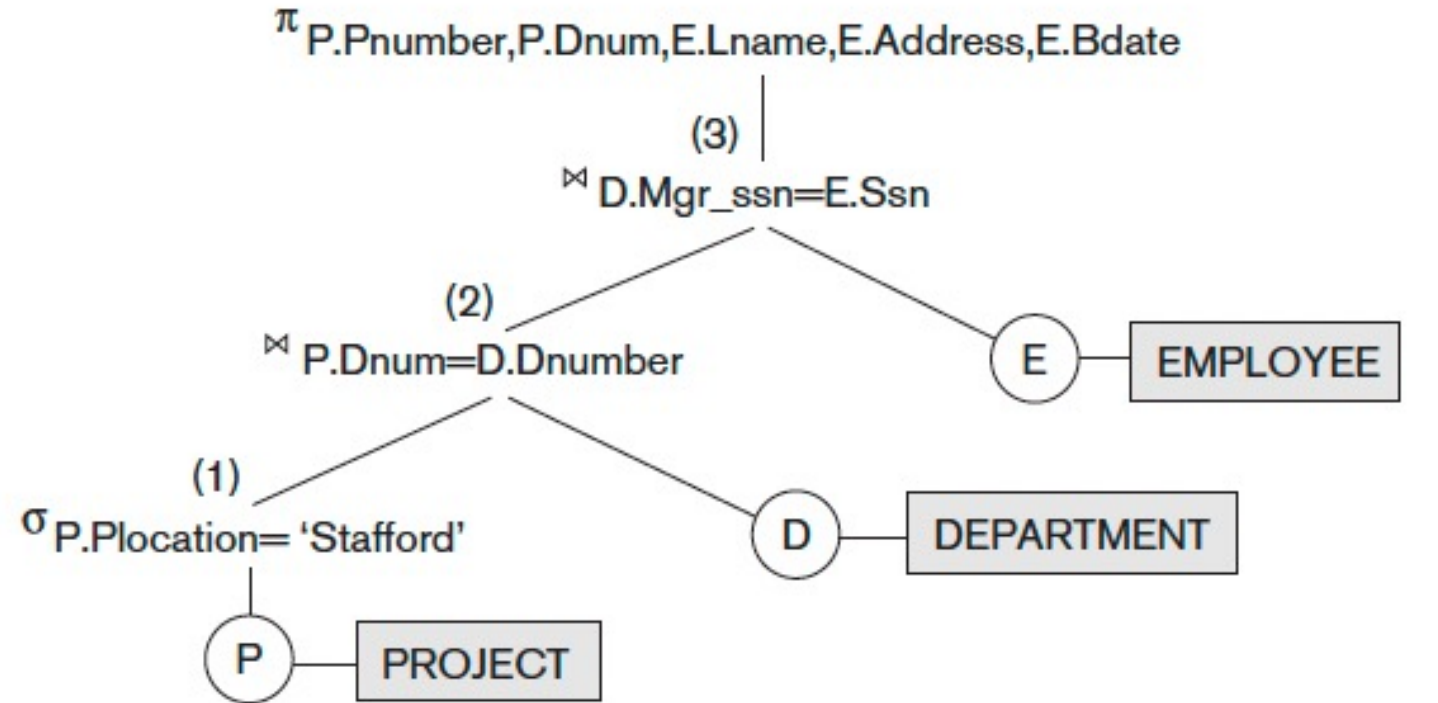
- The complete set of relational algebra operations  $\{\sigma, \pi, \cup, \rho, -, \times\}$ 
  - Select operation  $\sigma$
  - Project operation  $\pi$
  - Union operation  $\cup$
  - Rename operation  $\rho$
  - Set difference  $-$
  - Cartesian Product  $\times$
- \*The division operation  $\div$

# Query Tree Examples

**SELECT** P.Pnumber,  
P.Dnum,  
E.Lname,  
E.Address,  
E.Bdate  
**FROM** PROJECT P,  
DEPARTMENT D,  
EMPLOYEE E  
**WHERE** P.Dnum=D.Dnumber AND  
D.Mgr\_ssn=E.Ssn AND  
P.Plocation= 'Stafford';

(a) Query tree corresponding to the relational algebra expression

(a)



$\pi_{Pnumber, Dnum, Lname, Address, Bdate} (((\sigma_{Plocation='Stafford'}(PROJECT))$   
 $\bowtie Dnum=Dnumber(DEPARTMENT)) \bowtie Mgr\_ssn=Ssn(EMPLOYEE))$

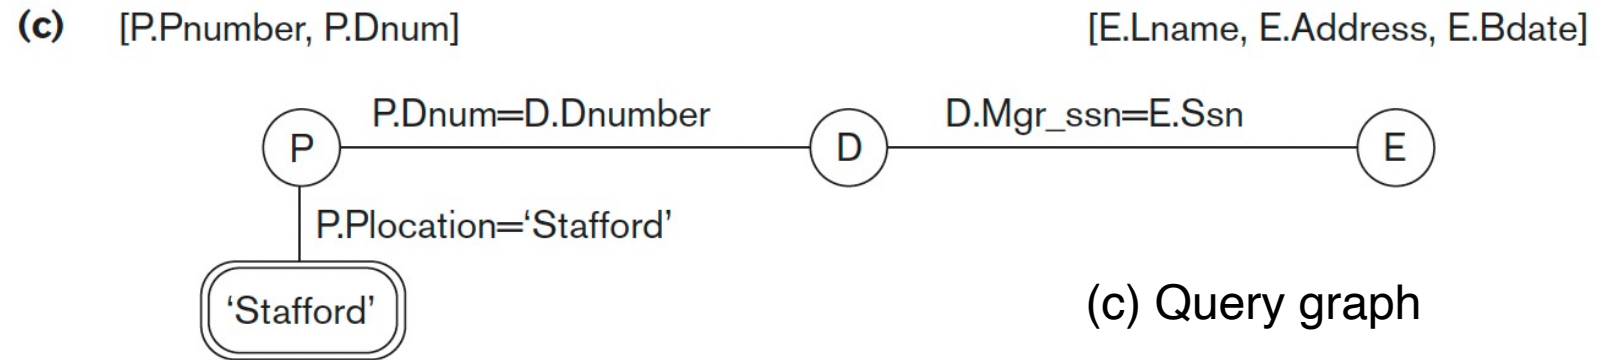
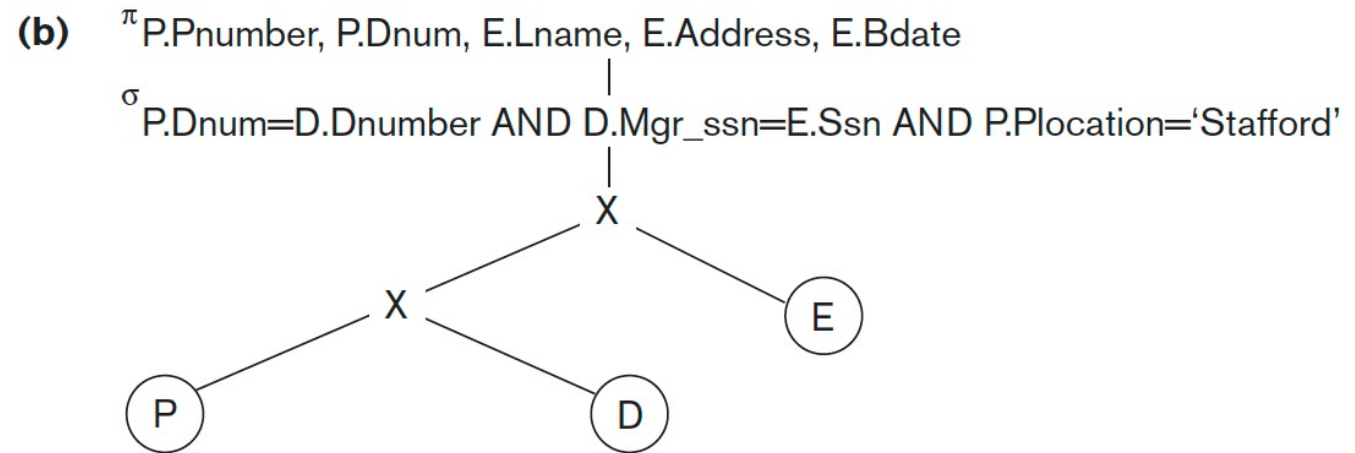


# Query Tree Examples

```

SELECT P.Pnumber,
       P.Dnum,
       E.Lname,
       E.Address,
       E.Bdate
FROM   PROJECT P,
       DEPARTMENT D,
       EMPLOYEE E
WHERE  P.Dnum=D.Dnumber AND
       D.Mgr_ssn=E.Ssn AND
       P.Plocation= 'Stafford';
    
```

(b) Initial (canonical) query tree for the SQL query



$\pi_{Pnumber, Dnum, Lname, Address, Bdate}(((\sigma_{Plocation='Stafford'}(PROJECT)))$   
 $\bowtie Dnum=Dnumber(DEPARTMENT)) \bowtie Mgr\_ssn=Ssn(EMPLOYEE))$

# General Transformation Rules for Relational Algebra Operations

- Cascade of  $\sigma$ . A conjunctive selection condition can be broken up into a cascade (that is, a sequence) of individual  $\sigma$  operations:

$$\sigma_{c_1 \text{ AND } c_2 \text{ AND } \dots \text{ AND } c_n}(R) \equiv \sigma_{c_1} (\sigma_{c_2} (\dots (\sigma_{c_n}(R)) \dots))$$

- Perform filtering as early as possible
- Break a complex predicate and push them down

# General Transformation Rules for Relational Algebra Operations

- Joins  $\bowtie$
- Commutative and associative

$$R \bowtie_c S \equiv S \bowtie_c R$$

$$(R \theta S) \theta T \equiv R \theta (S \theta T) \quad \theta \in \{\cup, \cap, \times, \bowtie\}$$


- The number of joins orderings for an n-way join is  $4^n$ 
  - Exhaustively enumerating all data is too slow...

# General Transformation Rules for Relational Algebra Operations

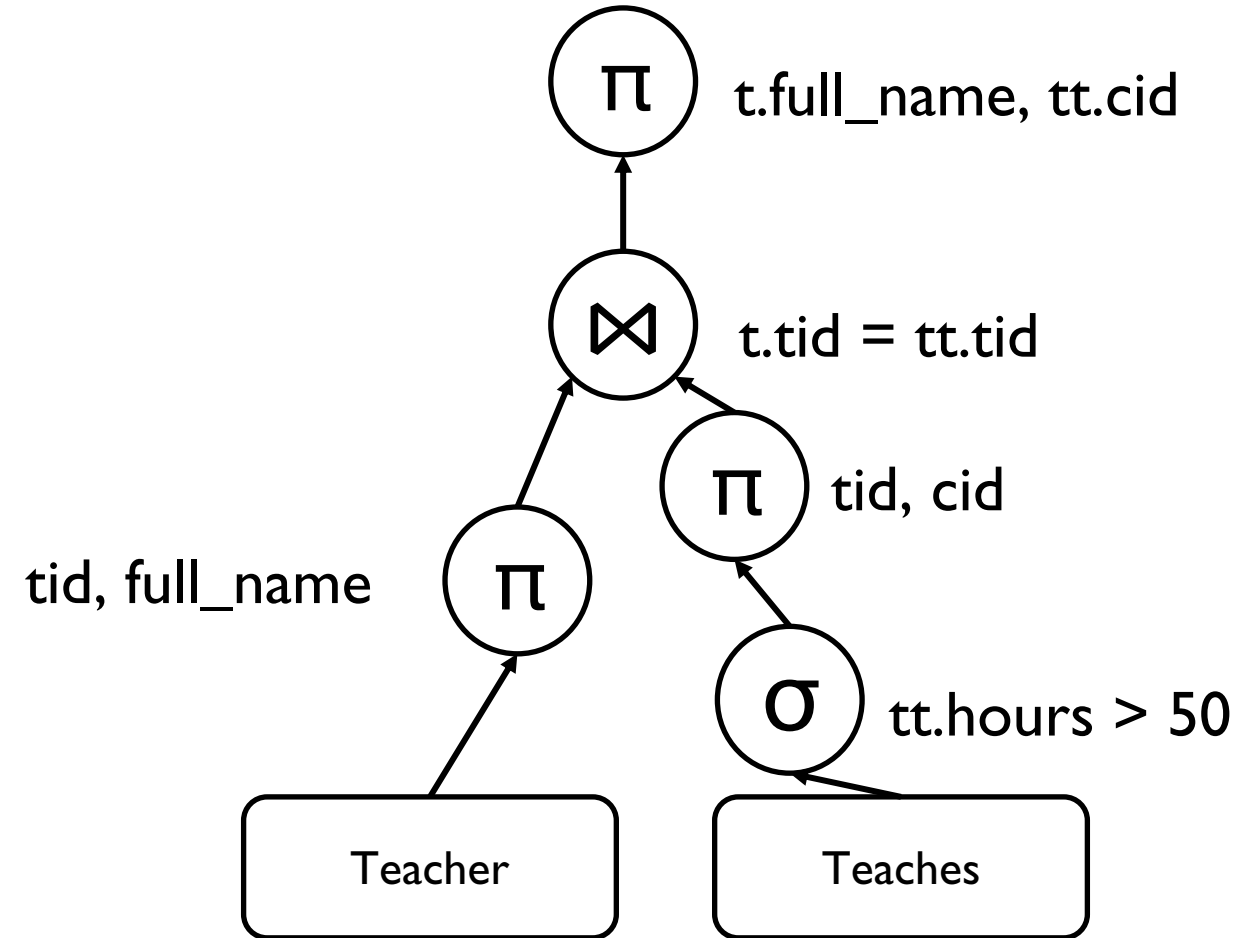
- Projections
  - Perform projects early to create smaller tuples and reduce intermediate results
  - Only takes the attributes that are required (e.g. joining keys) or requested by the query

# Example – Projection Pushdown

```
SELECT t.full_name, tt.cid  
FROM Teacher AS t, Teaches AS tt  
WHERE t.tid = tt.tid  
      AND tt.hours > 50
```



$\pi_{full\_name, cid}(\sigma_{hours>50}(Teacher \bowtie Teaches))$



# General Transformation Rules for Relational Algebra Operations

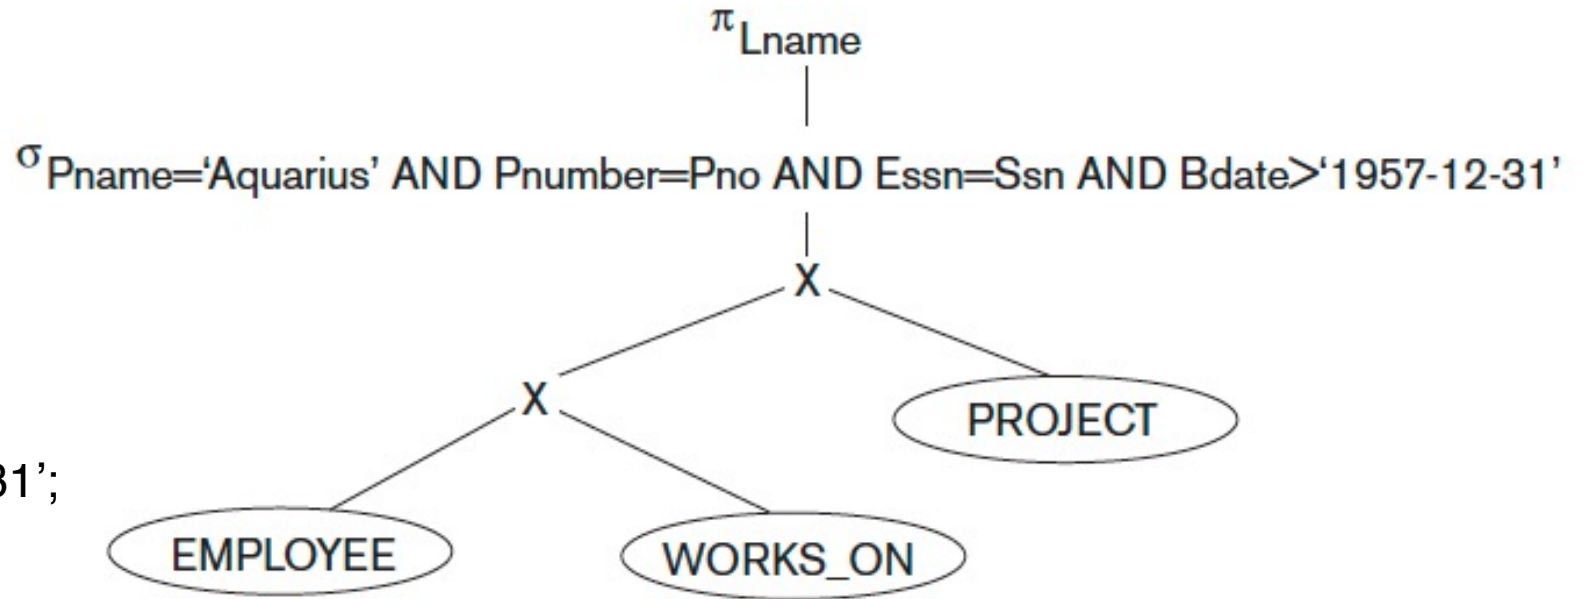
- Textbook page 698 – 699
- Cascade of selection and projection
- Commutativity of Selection
- Commuting selection with projection
- Commutativity of joins
- Commuting selection with joins
- Commuting projection with joins
- Commutativity of set operations
- Associativity of joins,  $\times$ , union and intersect
- ...

# General Transformation Rules for Relational Algebra Operations

- Outline of a Heuristic Algebraic Optimization Algorithm
  - Using Rule 1, break up any SELECT operations with conjunctive conditions into a cascade of SELECT operations.
  - Using Rules 2, 4, 6, and 10, 13, 14 concerning the commutativity of SELECT with other operations, move each SELECT operation as far down the query tree as is permitted by the attributes involved in the select condition.
  - Using Rules 5 and 9 concerning commutativity and associativity of binary operations, rearrange the leaf nodes of the tree using the following criteria.
  - Using Rule 12, combine a CARTESIAN PRODUCT operation with a subsequent SELECT operation in the tree into a JOIN operation, if the condition represents a join condition.
  - Using Rules 3, 4, 7, and 11 concerning the cascading of PROJECT and the commuting of PROJECT with other operations, break down and move lists of projection attributes down the tree as far as possible by creating new PROJECT operations as needed.

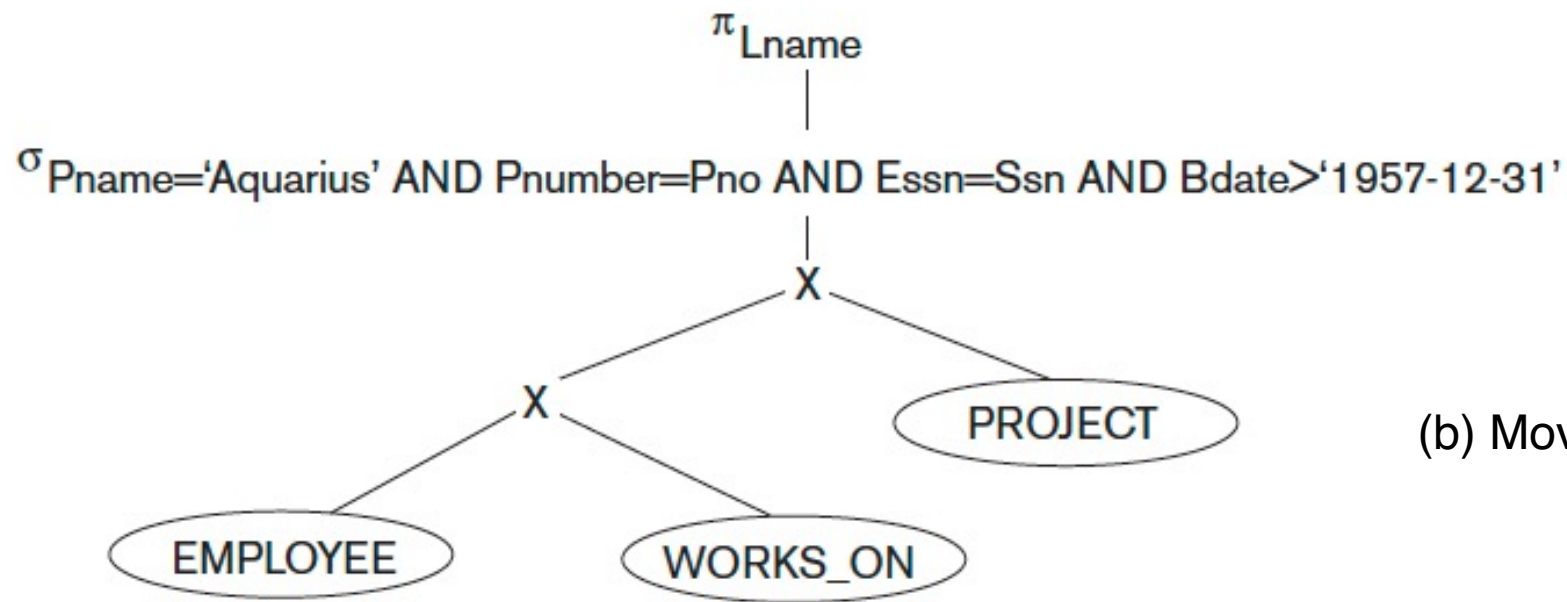
# Example transforming a query

```
SELECT E.Lname
FROM   EMPLOYEE E,
       WORKS_ON W,
       PROJECT P
WHERE  P.Pname='Aquarius'
      AND P.Pnumber=W.Pno
      AND E.Essn=W.Ssn
      AND E.Bdate > '1957-12-31';
```



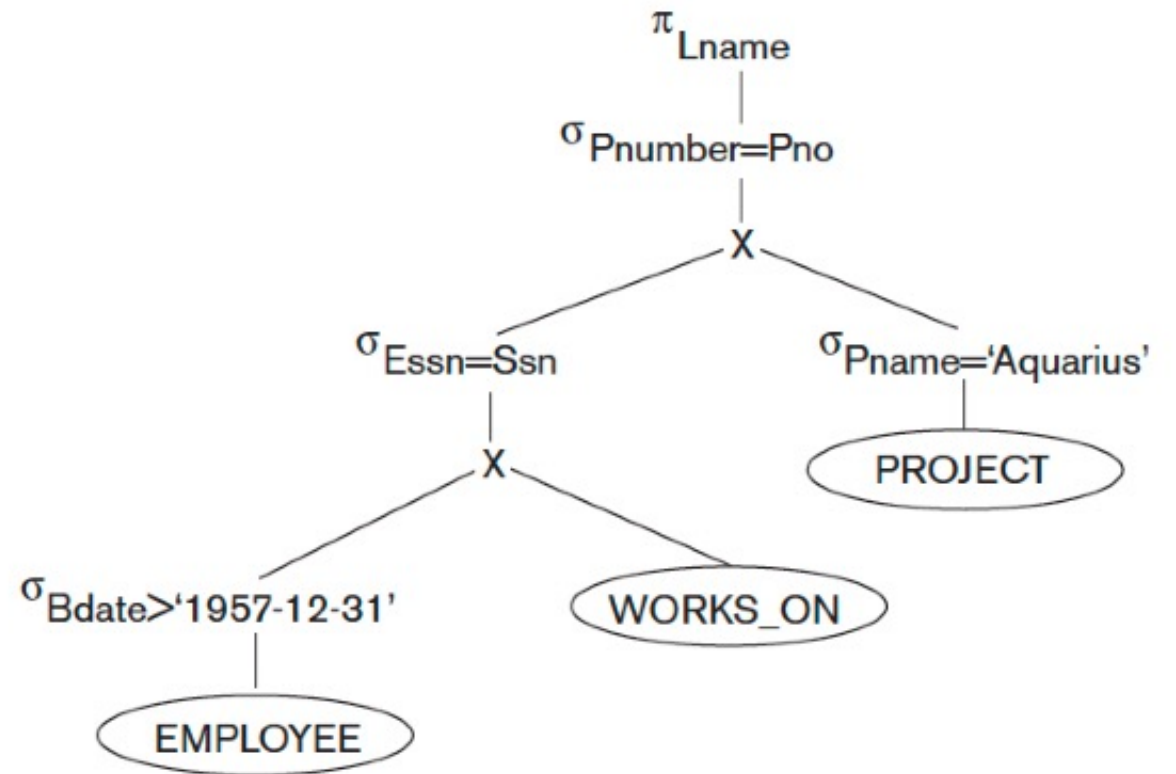
(a) Initial (canonical) query tree for SQL query

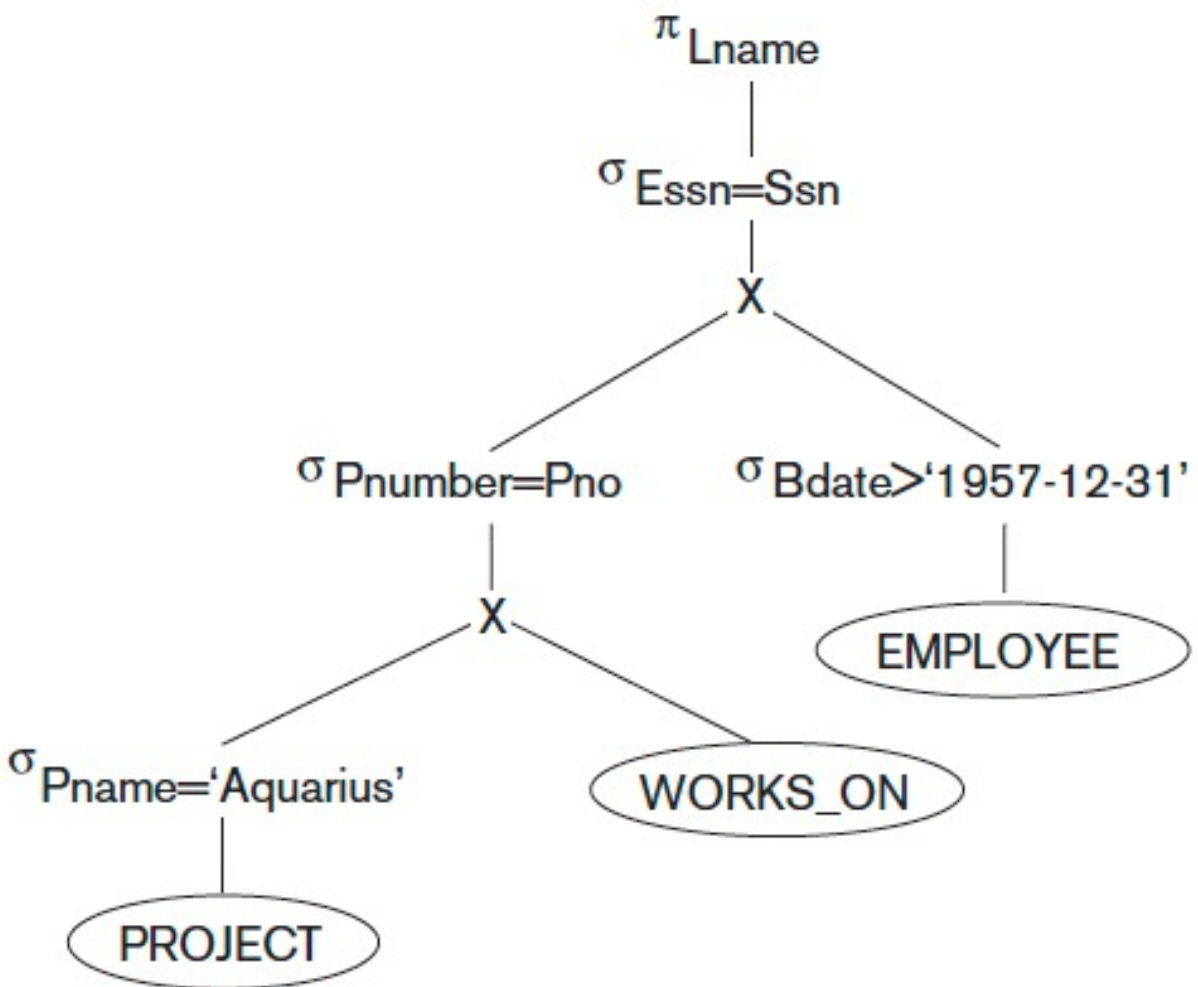




(a) Initial (canonical) query tree for SQL query

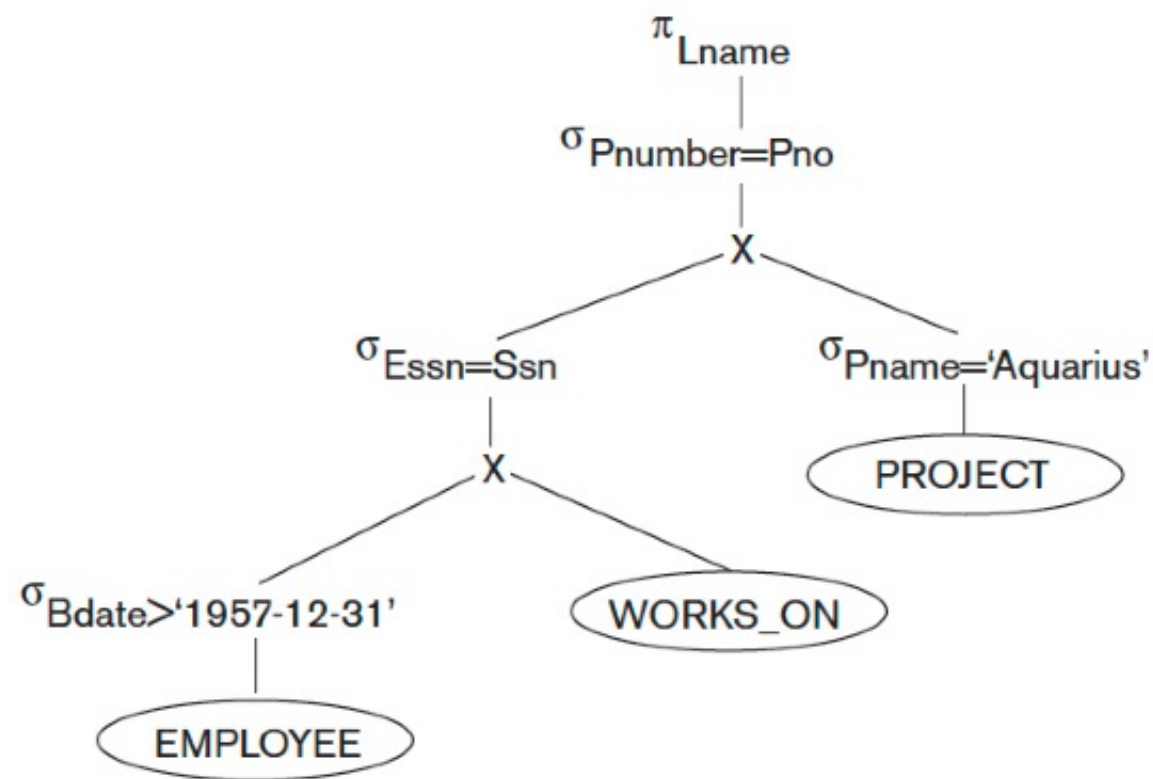
(b) Moving SELECT operations down the query tree.






(c) Applying the more restrictive SELECT operation first

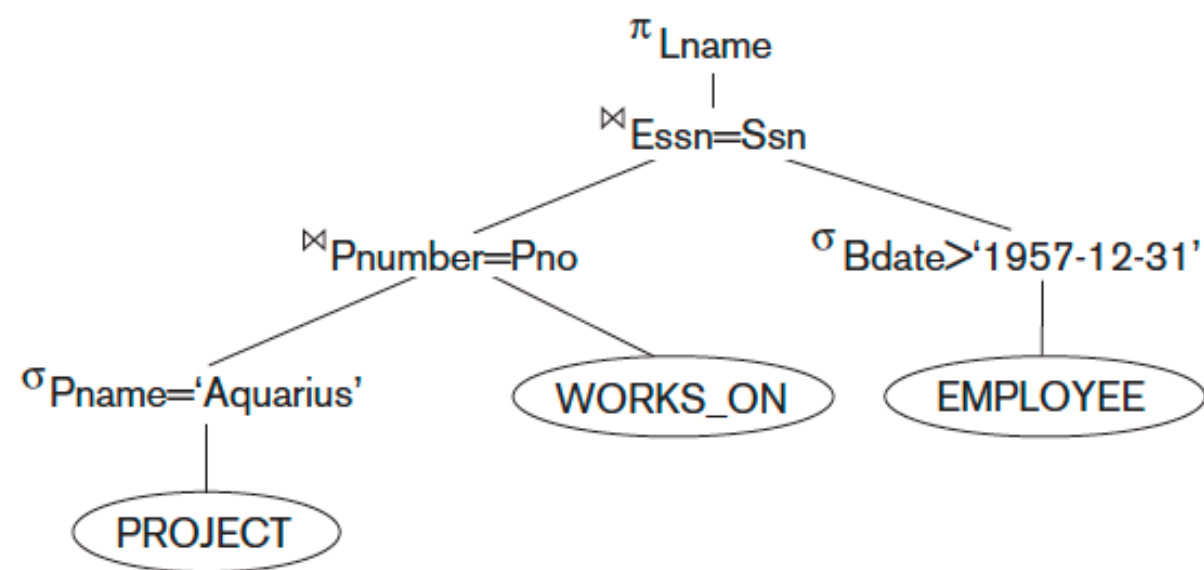
(b) Moving SELECT operations down the query tree.



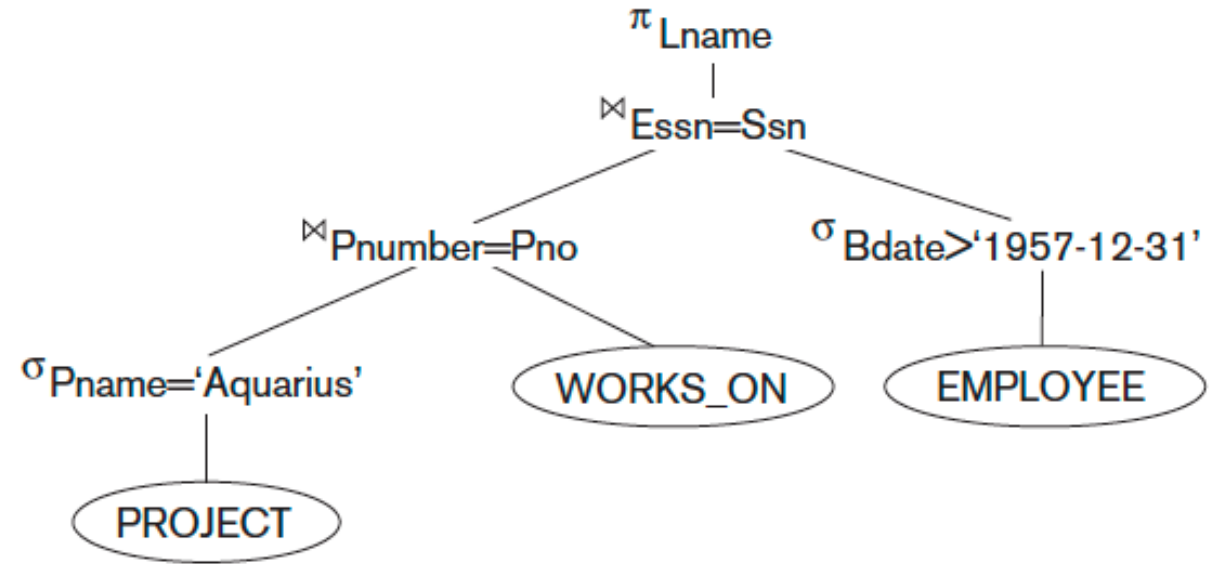
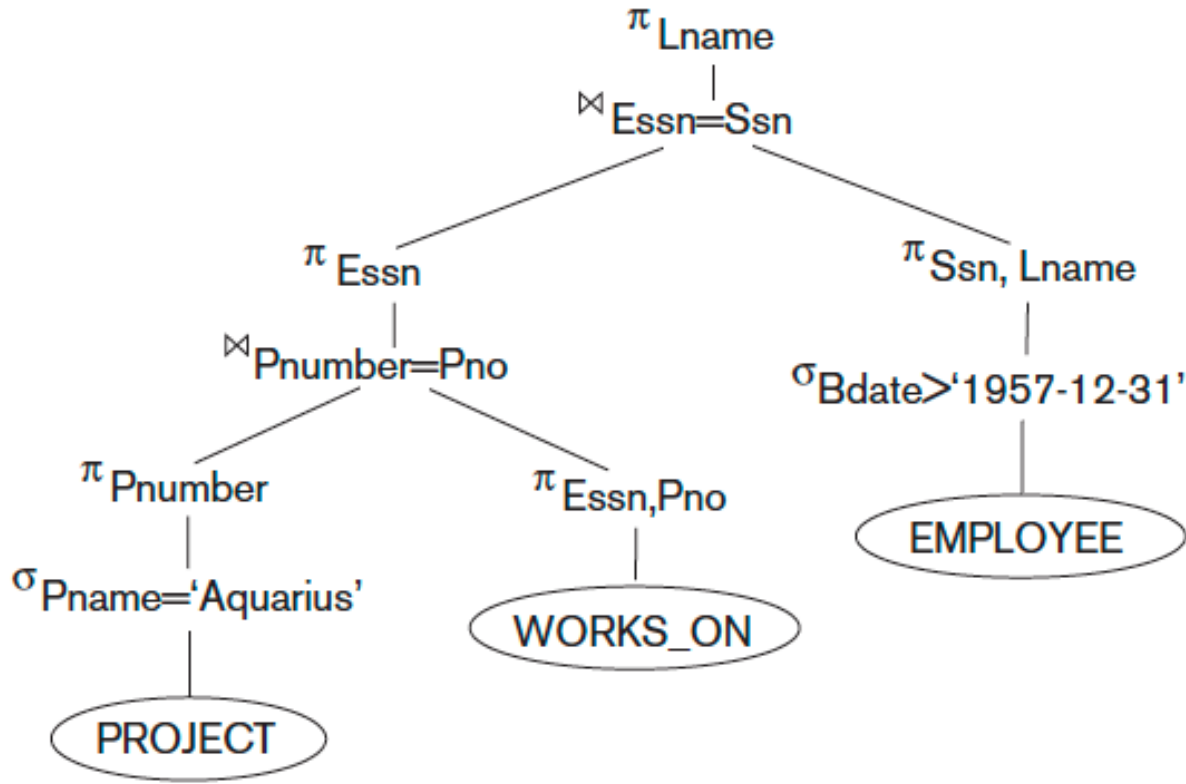


(c) Applying the more restrictive SELECT operation first

(d) Replacing CARTESIAN PRODUCT and SELECT with JOIN operations




(d) Replacing CARTESIAN PRODUCT and SELECT with JOIN operations



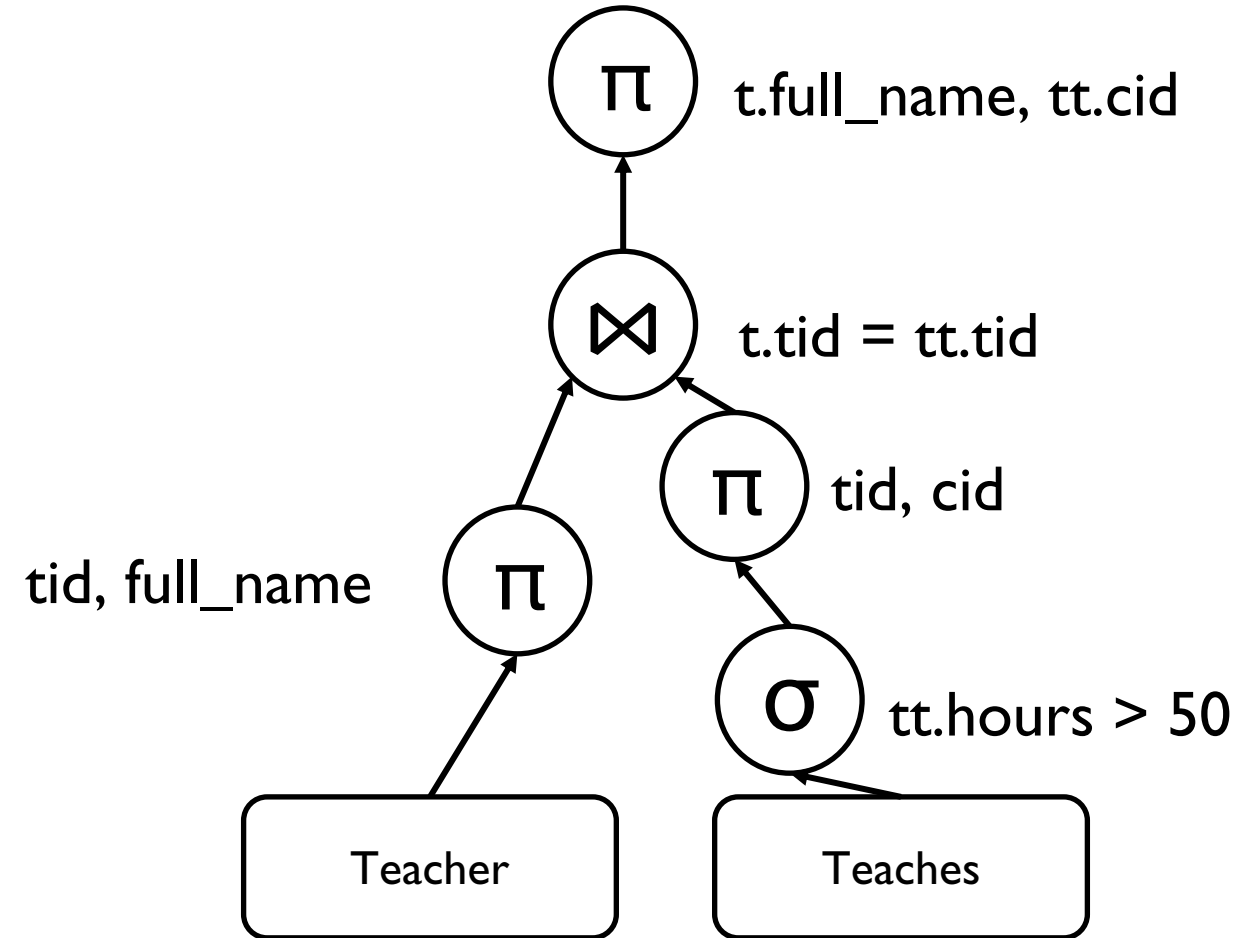
(e) Moving PROJECT operations down the query tree.

# Example – Projection Pushdown

```
SELECT t.full_name, tt.cid  
FROM Teacher AS t, Teaches AS tt  
WHERE t.tid = tt.tid  
      AND tt.hours > 50
```

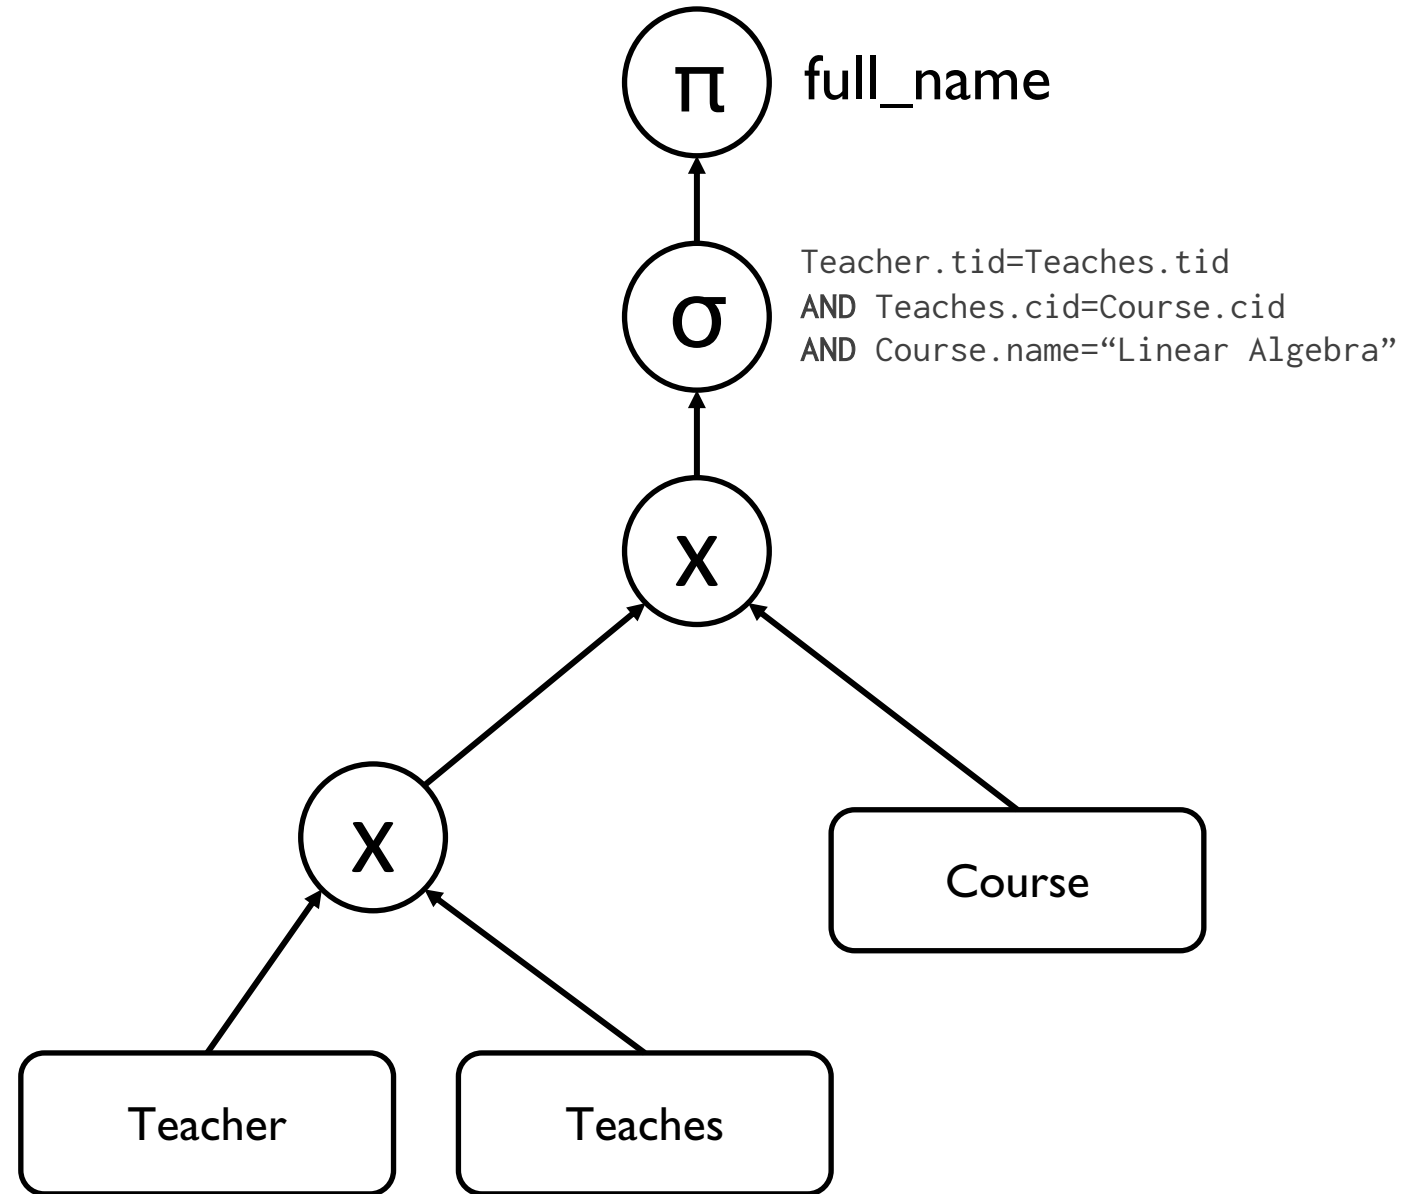


$\pi_{full\_name, cid}(\sigma_{hours>50}(Teacher \bowtie Teaches))$



# Example

```
SELECT Teacher.full_name
FROM Teacher, Teaches, Course
WHERE Teacher.tid=Teaches.tid
      AND Teaches.cid=Course.cid
      AND Course.name="Linear Algebra"
```

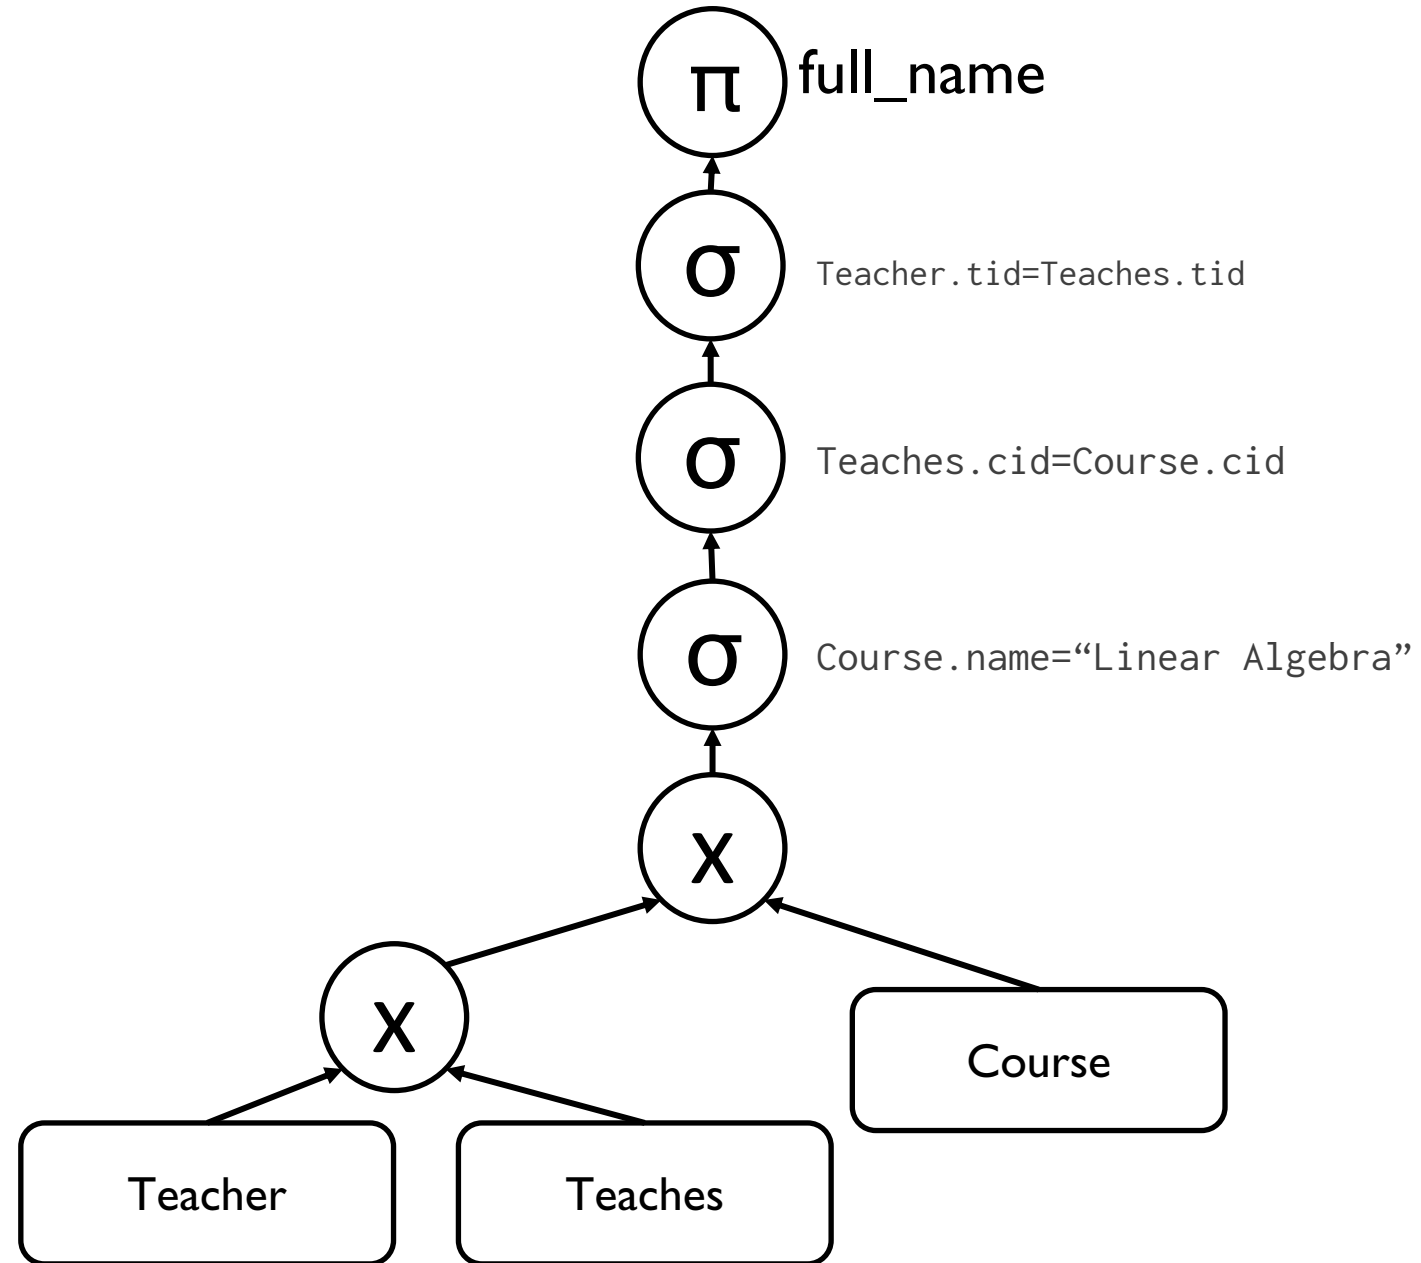


$\pi_{\text{full\_name}}(\sigma_{\text{t.tid=tt.tid AND t.cid=c.cid AND c.name="Linear Algebra"}}(\text{Teacher} \bowtie \text{Teaches} \bowtie \text{Course}))$

# Example

```
SELECT Teacher.full_name  
FROM Teacher, Teaches, Course  
WHERE Teacher.tid=Teaches.tid  
      AND Teaches.cid=Course.cid  
      AND Course.name="Linear Algebra"
```

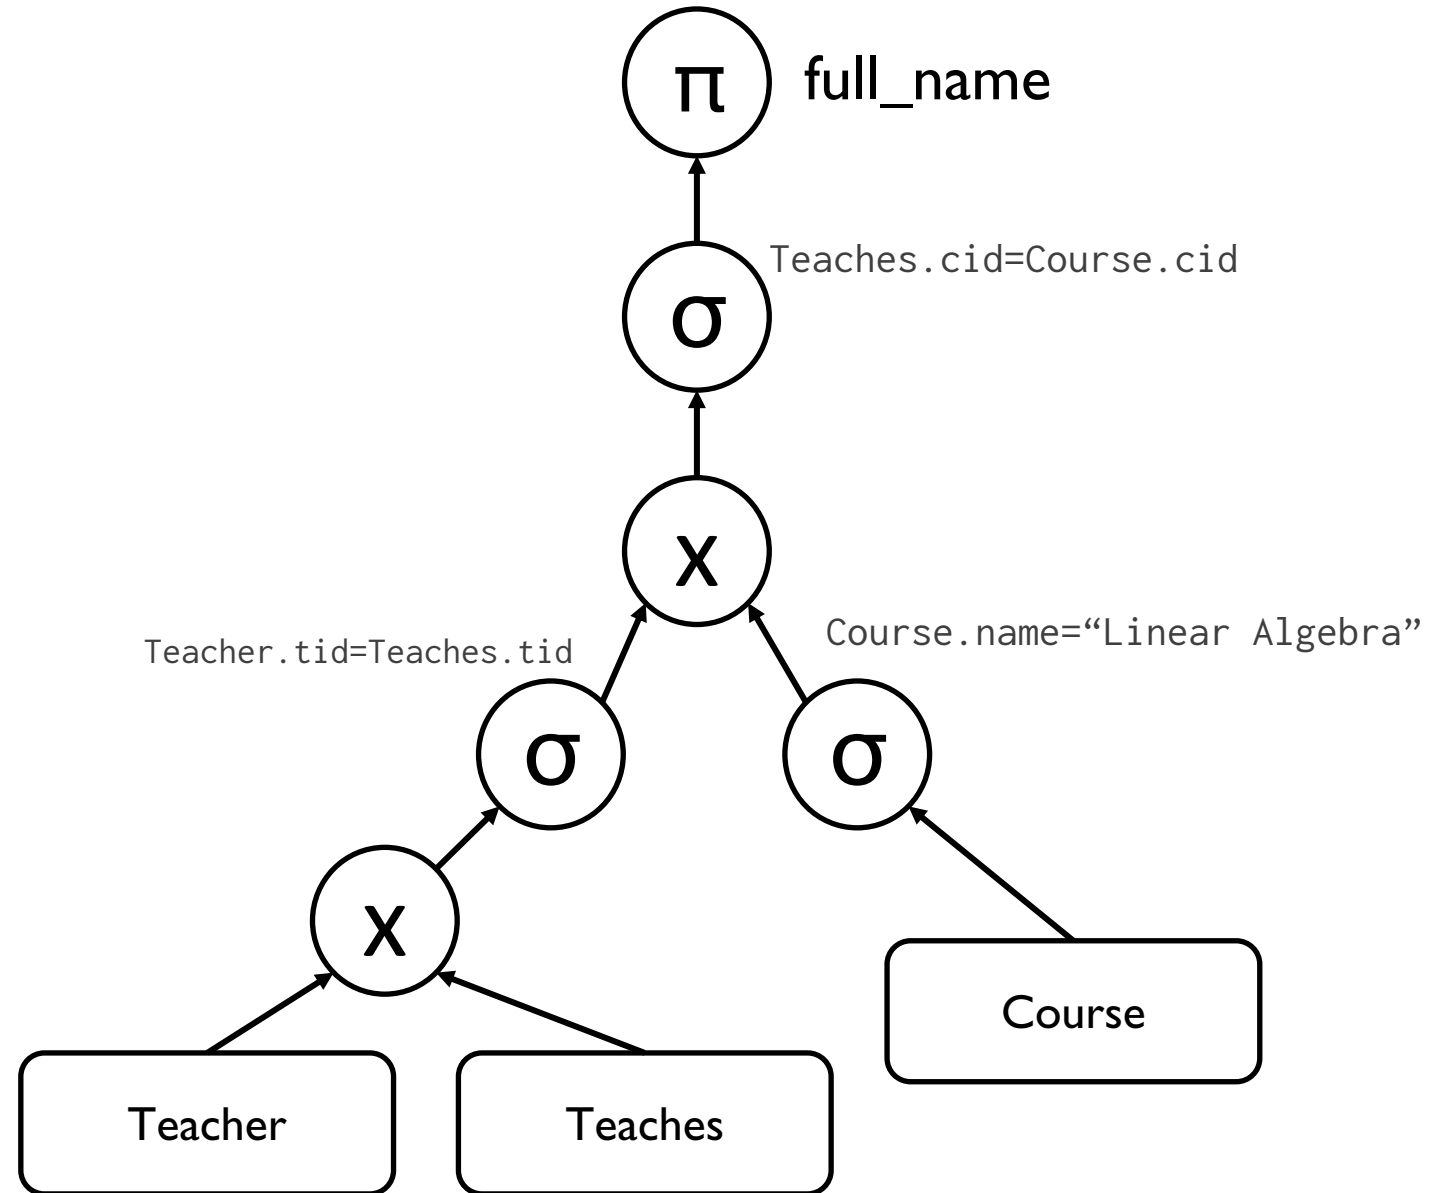
Decompose predicates into their simplest forms to make it easier for the optimizer to move them around.



# Example

```
SELECT Teacher.full_name
FROM Teacher, Teaches, Course
WHERE Teacher.tid=Teaches.tid
      AND Teaches.cid=Course.cid
      AND Course.name="Linear Algebra"
```

Move the predicate to the lowest point in the plan after Cartesian products.

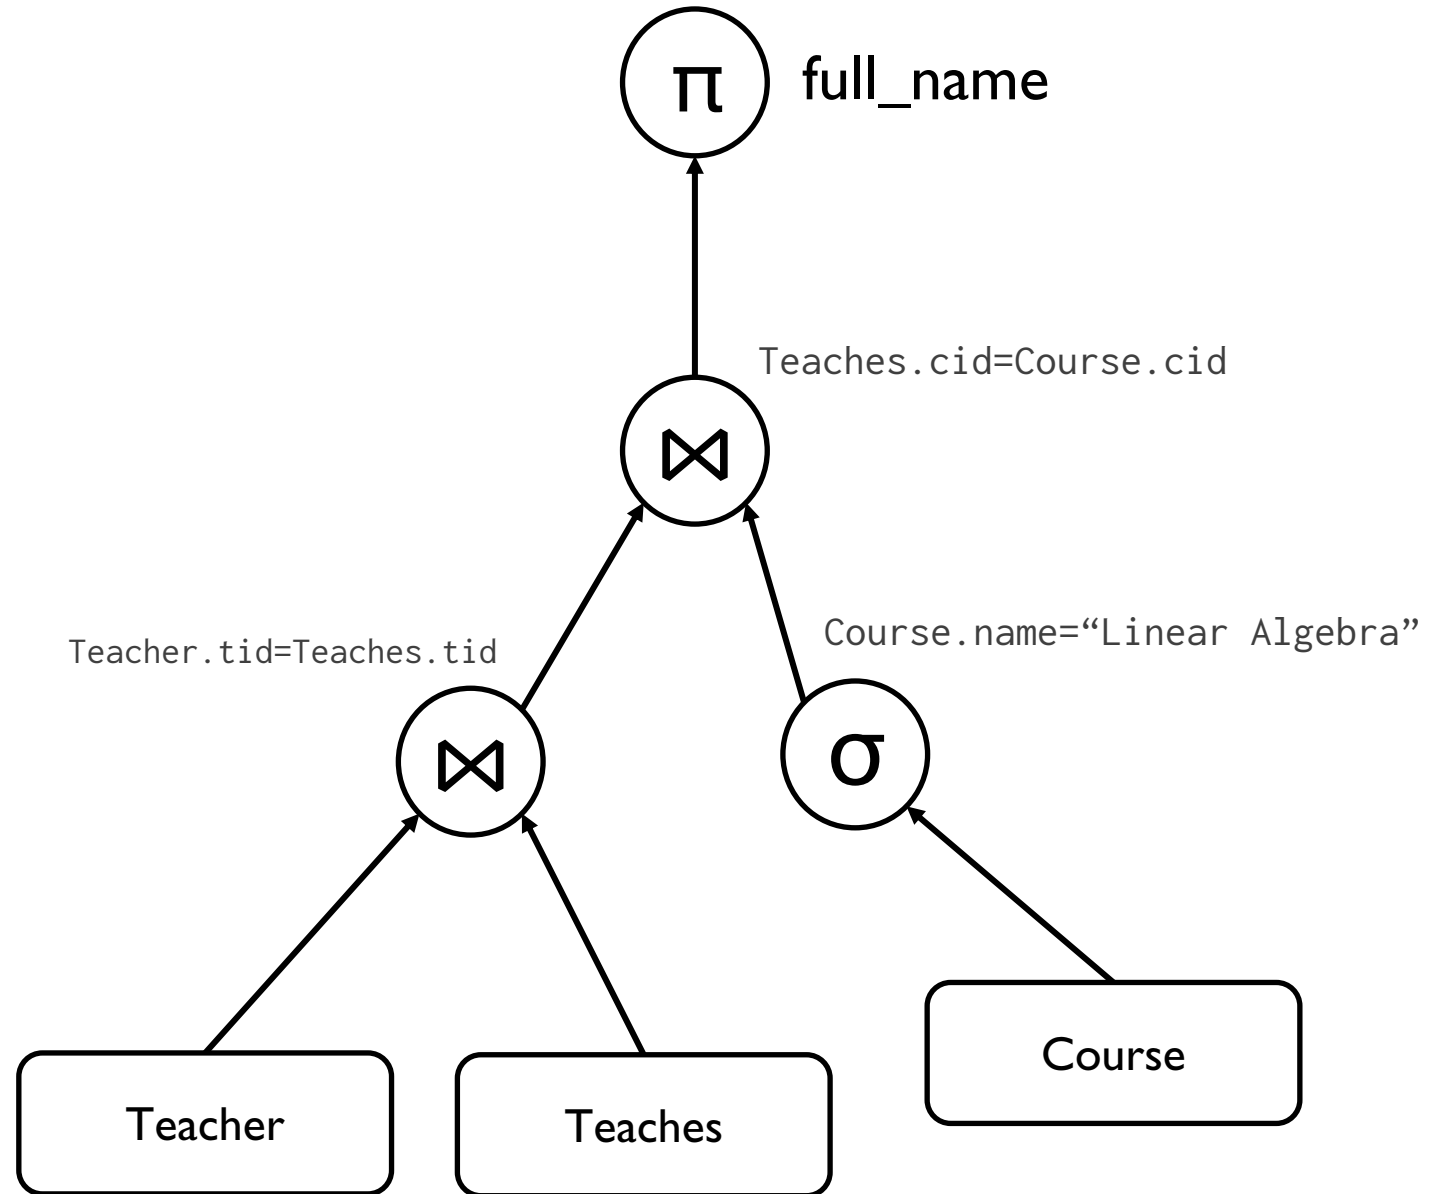




# Example

```
SELECT Teacher.full_name
FROM Teacher, Teaches, Course
WHERE Teacher.tid=Teaches.tid
      AND Teaches.cid=Course.cid
      AND Course.name="Linear Algebra"
```

Replace all Cartesian Products with inner joins using the join predicates.



# Example

```
SELECT Teacher.full_name
FROM Teacher, Teaches, Course
WHERE Teacher.tid=Teaches.tid
      AND Teaches.cid=Course.cid
      AND Course.name="Linear Algebra"
```

Eliminate redundant attributes before  
pipeline breakers to reduce materialization  
cost.

